

What an AI-friendly file actually looks like

Same content, two outcomes. The difference is a descriptive filename and a small block of metadata at the top. Here is an annotated example of a policy document built so an AI assistant can find it, trust it, and answer from it accurately.

THE IDEAL FILE

hr-policy-parental-leave-us-2026-v3.md

```
---
title: Parental Leave Policy (US)
doc_type: hr-policy
owner: People & Culture
audience: all-us-employees
topics: [parental-leave, benefits, time-off]
effective_date: 2026-01-01
last_reviewed: 2026-05-14
version: 3.0
status: active
supersedes: v2.1 (2024)
---
```

Parental Leave Policy (United States)

Effective Jan 1, 2026 · Reviewed May 14, 2026 · Version 3.0 · Owner: People & Culture

Summary. Eligible US employees receive paid parental leave following the birth, adoption, or foster placement of a child.

Who is eligible

- Full-time US employees, active 90+ days
- Birth, adoptive, and foster parents

What you get

- 12 weeks paid leave at 100% of base pay
- Leave may be taken within 12 months of the event

How to request

- Submit in the HR portal at least 30 days in advance

WHY EACH PIECE HELPS THE AI

1 Descriptive filename

The name encodes type, topic, region, year, and version. The AI can match the right file *before it even opens it*, no guessing from `policy_final.md`.

2 Metadata header (front matter)

The block between the --- lines is structured "data about the data." It sits in a predictable spot the AI reads first.

3 Topics & audience

Keywords like *parental-leave*, *benefits* are exactly what employees type. They become the AI's matching hooks.

4 Dates, version & status

last_reviewed, *version*, *status: active*, *supersedes* tell the AI this is current and which old copy to ignore, so it never answers from a stale policy.

Tidy for humans ≠ ready for AI

Humans navigate by **location** (folders). AI navigates by **name and content**. You don't need this on every file. Focus it on the static, high-traffic stuff people ask AI about: policies, SOPs, benefits, and knowledge-base articles.

● THE FILENAME, BEFORE AND AFTER

HARD FOR AI

policy_FINAL_v2.docx
updated-doc(1).pdf
leave.md

EASY FOR AI

hr-policy-parental-leave-us-2026-v3.md
sop-expense-reimbursement-2026.md
benefits-401k-match-us-2026.md

1

Name it like a search query

Encode type, topic, region, and year in the filename. It's the highest-value, lowest-effort metadata there is.

3

Signal what's current

Mark *status*: *active*, a review date, and what it supersedes. This is how the AI avoids answering from an outdated version.

2

Add a metadata header

A short block at the top: title, owner, topics, dates, version, status. It lives *inside* the file, readable by both people and AI.

4

Do it where it counts

Static, high-traffic docs in SharePoint, ServiceNow, or a knowledge base. Skip it for one-off and fast-changing files.

THE TAKEAWAY

When someone asks an AI assistant "how much parental leave do I get?", it finds the answer instantly and confidently when the file is **named like the question** and **carries metadata that proves it's current**. A little structure today is a far better answer tomorrow.